

## Training over three stages



Loss values are normalized to each stage's first epoch: the three stages optimize different quantities and their absolute values are incomparable. RL reward is shown only for stage C.

Figure 4 — training curves across three stages and RL reward